

Data Mining & Machine Learning F20DL

Coursework 2

<https://github.com/friedforfun/SupervisedLearningCW>

Group 4

Lewis Wilson, Sam Fay-Hunt, Kamil Szymczak, Chun Man

December 8, 2020

Contents

1	Introduction	1
2	Variation in performance with size of the training and testing sets	2
3	Variation in performance with the change in the learning paradigm	5
3.1	Convolutional Neural Networks	6
3.1.1	Critical limitation of CNN experimental methodology:	6
3.1.2	Summary of CNN experimental results	6
4	Methodology for Optimising hyperparameters	8
5	Variation in performance with varying learning parameters in Decision Trees	9
5.1	J48	9
5.2	Random Forest	10
6	Variation in performance with varying learning parameters in Neural Networks	11
6.1	Linear Classifier	11
6.2	Multilayer Perceptron	12
6.3	Further experiments	12
6.3.1	Turn left label	13
6.3.2	Beware Pedestrian Crossing - Looking closer	13
6.3.3	Conclusions about the extended MLP sweep	14
7	Variation in performance according to different metrics	15
7.1	J48 Experiments	15
7.2	Random Forest & MLP	15
7.3	Logistic Regression	15
A	Appendices	16
A	Workload	16

A.1 Workload split	16
B References	17
C Links to experimental results	18
C.1 J48	18
C.2 Random Forest	18
C.3 MLP	18
C.4 LR	18
C.5 CNN	19
D J48	20
D.1 J48 Parameter Importance	20
E Random Forest	21
E.1 Random Forest Parameter Importance	21
F Linear Classifier	22
F.1 Linear Classifier Parameter Importance	22
G Multilayer Perceptron	23
G.1 Multilayer Perceptron Parameter Importance	23

1 Introduction

Based on our experience from Coursework 1 we were aware that many machine learning models would perform poorly with a poorly balanced class distribution, we took measures to resolve this class distribution problem in coursework 1 and this resolved the overfitting issues. We chose not to use these techniques so we could test the machine learning models with an inadequate dataset, to give us insight into how the different models perform in this situation.

We wrote all experiments in Python 3, using SKlearn because it was easy to implement the models under an interface and pass the discrete models as parameters to “run experiment” functions, we used Tensorflow+Keras for the CNN implementation for ease of use, and speed of training (because of easy integration with Cuda). To record and visualise experimental data we used Weights and Biases, this allowed us to perform experiments discretely and share all the results as they were generated.[1]

The purpose of these experiments is to explore how variances in data, learning paradigm and hyperparameters can affect machine learning model performance with respect to overfitting. To do this we ran 4 Experiments with each model, covering two machine learning paradigms, we could then compare the behaviour, and try to gain insight into if and why they perform so poorly with data using bad class distributions. For all but the CNN we recorded the Accuracy, F1, Precision, and Recall metrics, additionally we plotted the confusion matrix, precision-recall curves, ROC curve, class proportions, calibration curve, and learning curve. [1]

The 4 experiments that were run are as follows:

- K-fold - We used stratified K-fold because the class label distribution is so unbalanced each of our folds would not be guaranteed to have all classes represented in the training set, this would have unnecessarily skewed our results by testing on never-before-seen data, and would make drawing conclusions much harder. (Figure 1)
- Train-test - Used the provided datasets for training and for each label set a testing vector.
- 4000 and 9000 instances moved from train to test - Same as train-test split above but with the difference being 4000 and 9000 instances moved from train to test. This allowed us to evaluate the model’s performance with different proportions resulting in less training samples.

We used the same seed when moving the data from the training set to the testing set, so all experiments for a given classifier used the same training and testing data.

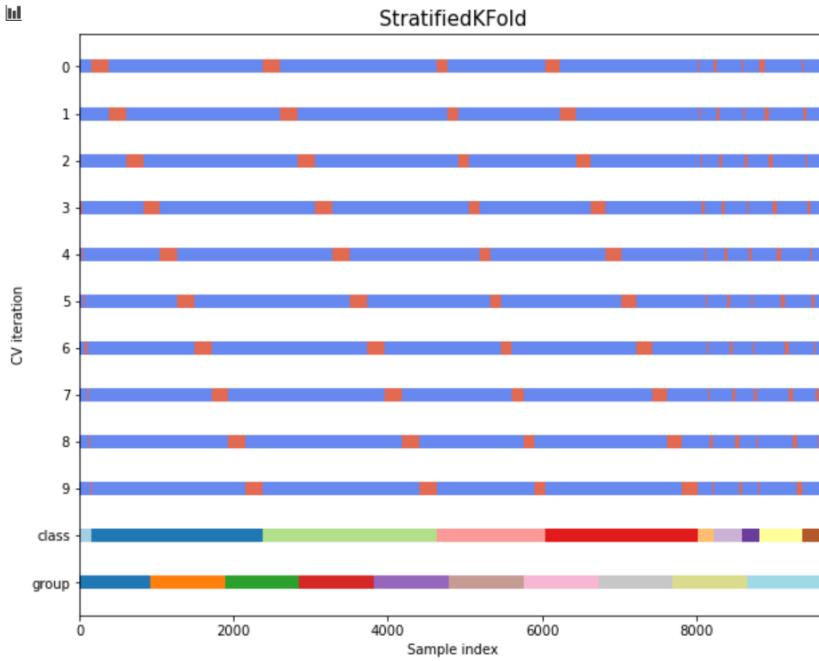


Figure 1: Stratified KFolds, Visualised class distributions in each fold (Group)

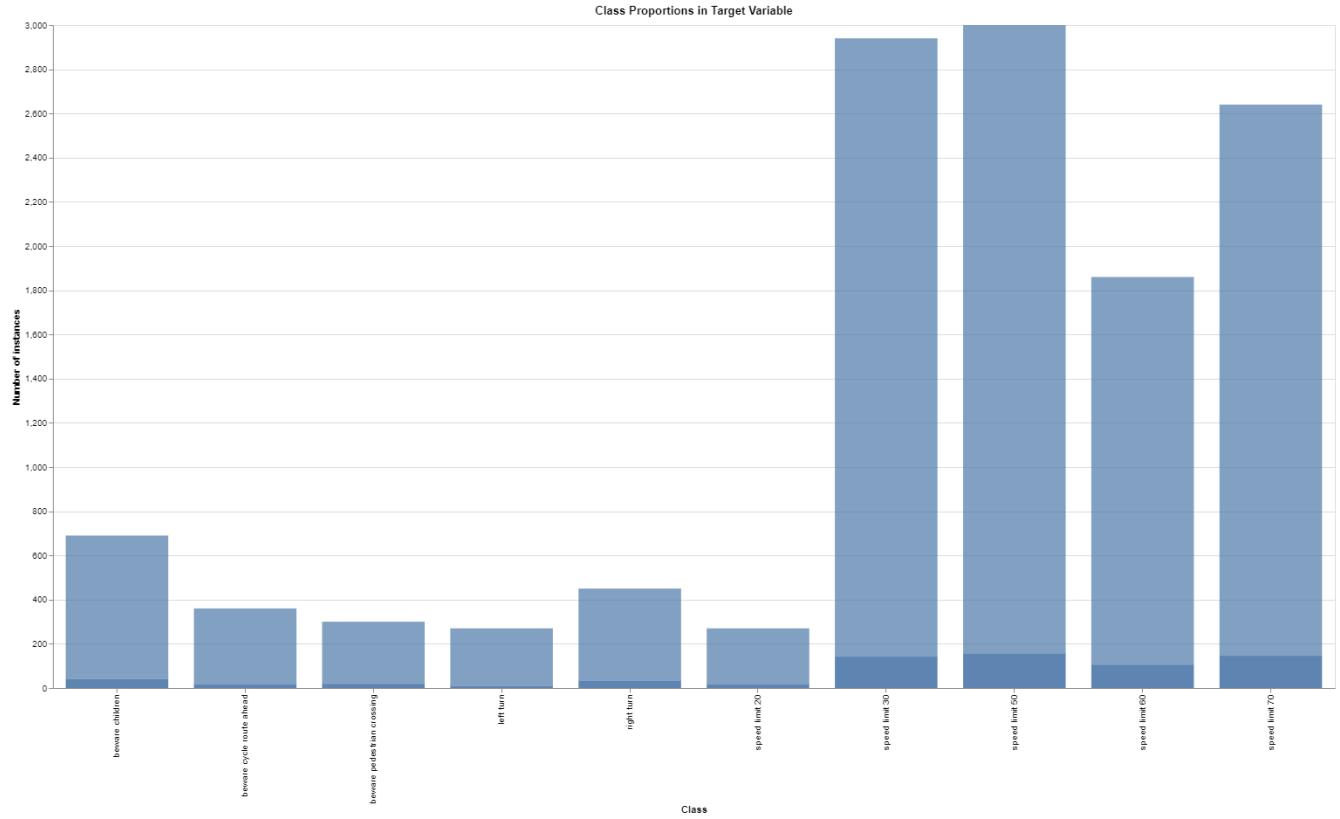
2 Variation in performance with size of the training and testing sets

To investigate the variation in performance we began by looking closely at a single class with a very poor class distribution: “Turn left”.

The turn left dataset has an extremely small number of instances in its training dataset for 9000 moved to testing, Thus we would expect the worst prediction accuracy, in this case, the situation was so bad that our accuracy scores became misleading. Because the representation of ‘turn left’ was so poor in our dataset most of the classifiers simply predicted that there were no instances of turn left in the dataset and recorded a deceiving 97% accuracy (because 97% of the dataset is not in fact turn left). It turned out this problem appeared for all binary classifications with the Random Forest classifier and Multilayer perceptron.

The only classifiers that did not exhibit this behaviour were the Logistic Regression (LR) and Convolutional Neural Network classifiers (CNN). CNN scored a staggering 99.9% accuracy with the same minuscule training data (without out enhancements to reduce overfitting, see Section x), while LR managed 98.3% with 200 false negatives, 5 false positives and 60 true positives.

Figure 2: All class proportions of train (dark blue) to test (light blue)



An example distribution of labels between train (dark blue) and test (light blue), for the Beware cycle route ahead class labels, Figures 3, 4, 5 show how the proportion of training data moves more substantially for the ‘no’ class, while the ‘yes’ class is so small the change is marginal.

Figure 3: Beware cycle route ahead class proportions

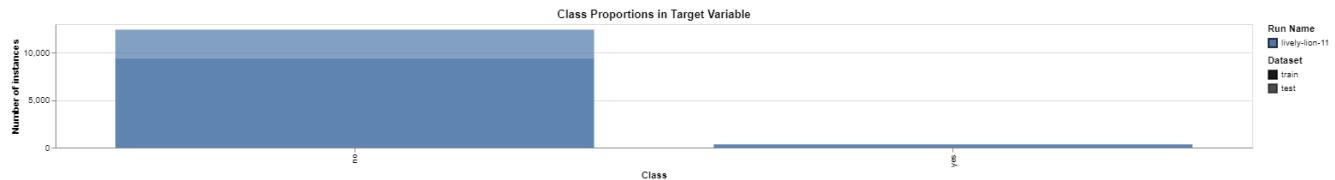


Figure 4: Beware cycle route ahead class proportions with 4000 moved to tests

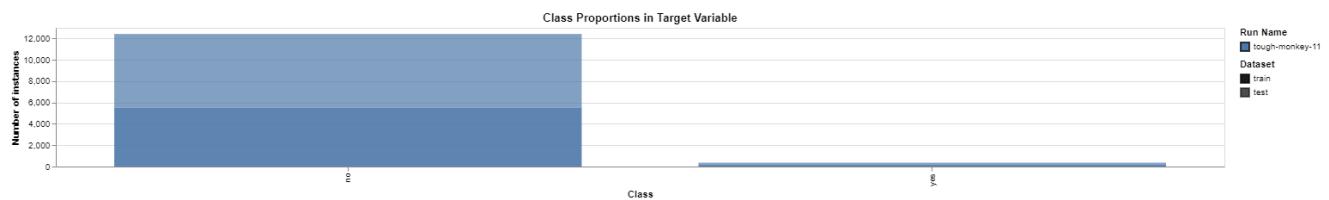
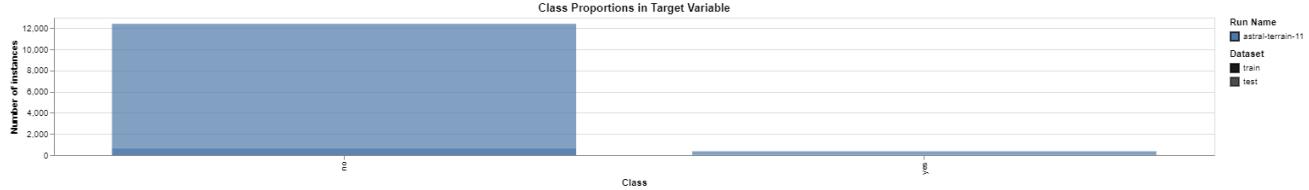
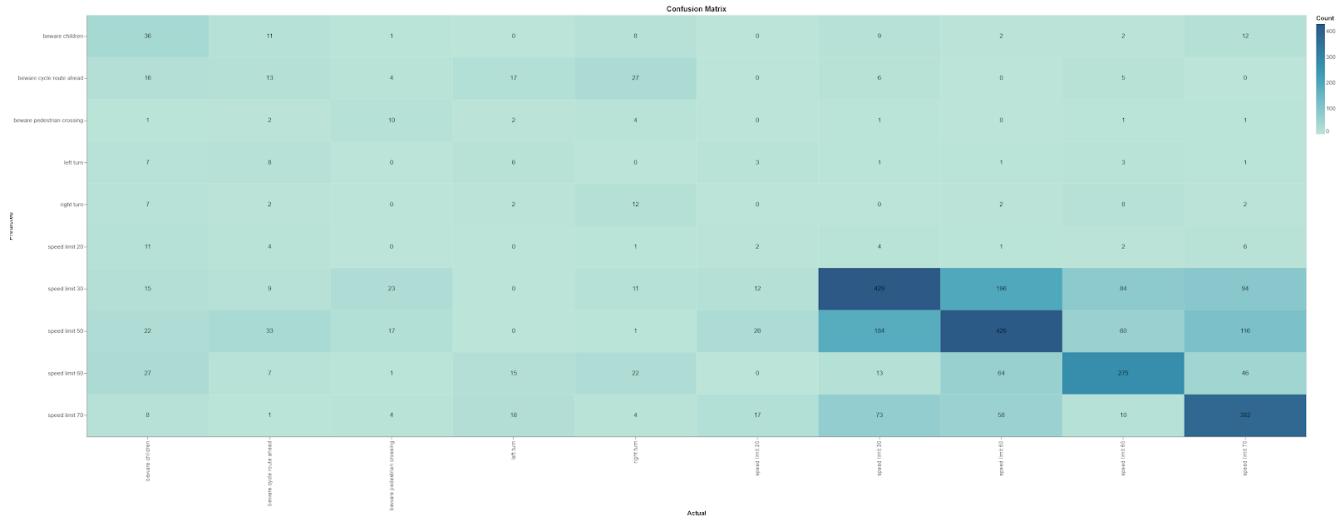


Figure 5: Beware cycle route ahead class proportions with 9000 moved to tests



Starting out, we observed the results of J48, the issues are immediately clear, even in the experiment with the full training set J48 is poor at attempting to predict on this kind of data. When working with ‘all classes’ it is only able to classify the very well represented speed limit signs reliably as speed limit signs, so it manages a high precision at classifying a meta-class of the group of speed limit 30, 50, 60, or 70 signs. Within this meta-class the accuracy at distinguishing between them is approximately 50% on average.

Figure 6: J48 Confusion Matrix

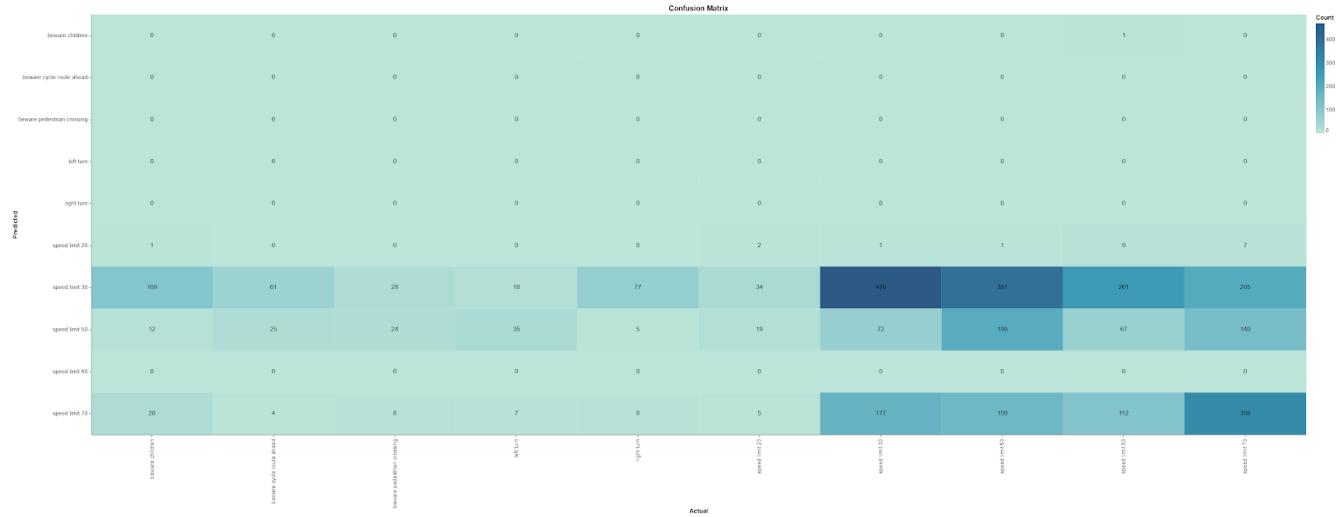


As the test set grows we can observe that the J48 classifier gets substantially worse at distinguishing between speed limit signs and other signs because the representation of those signs in the dataset gets very low.

With all the classifiers except Logistic Regression and the CNN we see the same story, Multi-Layer perceptron and random forest particularly suffer from this poor class representation issue in the data. We can see in Figure 7 that the Multi-Layer perceptron assigns most labels for 3 classes: Speed limit 30, 50, 70. It also assigns a small number of labels for speed limit 20. So in this baseline test it is already showing overfitting behaviour.

In the test where 9000 instances are moved the overfitting behaviour gets much worse, the Multi-Layer perceptron classifies everything as either speed limit 30 or speed limit 70.

Figure 7: Multilayer Perceptron Confusion Matrix



3 Variation in performance with the change in the learning paradigm

- Decision trees performed comparatively poorly, particularly when the training data was reduced.
- Multilayer Perceptron had overfitting issues, Logistic Regression was the best, aside from the CNN

We have found decision trees performed very poorly particularly when the training data was reduced and also so did the MultiLayer Perceptron which had overfitting issues. Both of these classifiers had 0 recall and 0 precision for predicting individual classes datasets. This means they have not predicted correctly a single sign and should not be used in prediction road signs from images. Our results perfectly support the weak points about decision trees such as that there are variable relationships that decision trees just can't learn and they are unstable as a small change in the data can lead to a large change in the structure of the optimal decision tree [2].

However, Logistic Regression performed considerably better, it was actually the best performer from all classification techniques we have used and tested, aside from the CNN.

Logistic Regression has given best results thus it was less vulnerable to poor class distributions and dealt nicely with overfitted data. A benefit of logistic regression is its interpretability as it understands which pixels are important to determinate what class an image belongs to.

Convolutional Neural Networks worked phenomenally as expected, as they are widely used in image classification.

3.1 Convolutional Neural Networks

Table 1: Shape of custom CNN

Layer type	Output Shape	Param #
conv2D	(None, 48, 48, 75)	750
BatchNormalization	(None, 48, 48, 75)	300
MaxPooling2D	(None, 24, 24, 75)	0
Conv2D	(None, 24, 24, 50)	33800
Dropout	(None, 24, 24, 50)	0
BatchNormalization	(None, 24, 24, 50)	200
MaxPooling2D	(None, 12, 12, 50)	0
Conv2D	(None, 12, 12, 25)	11275
BatchNormalization	(None, 12, 12, 25)	100
MaxPooling2D	(None, 6, 6, 25)	0
Flatten	(None, 900)	0
Dense	(None, 512)	461312
Dense	(None, 512)	262656
Dropout	(None, 512)	0
Dense	(None, 10)	5130

3.1.1 Critical limitation of CNN experimental methodology:

We were unable to disable memoization between experiments over each loop, so all the parameters are restored between runs within a given experiment, consequently, we noted only significant changes to the loss when training on all classes, after which convergence was achieved in a single epoch. We attempted to clear the model parameters, clear the GPU memory, and model saves the randomly initialised parameters and load them back in between experiments but could not prevent the CNN model using data encoded from previous experiments. This does invalidate our results when compared to the discrete experiments in the previous sections, however, this can be a demonstrable advantage of a CNN, in this situation it consistently correctly classifies the images despite changing the dimensions of the output layer and the classification targets (albeit on the same dataset).

3.1.2 Summary of CNN experimental results

During all experiments with training and testing data, at worst our CNN scored an accuracy of 97% on the test dataset when classifying all classes, the CNN was capable of correct classification even when 9000 instances were moved from the training set to the testing set.

The CNN was unsurprisingly the best performing classifier, In many of the k-fold experiments the CNN

managed an accuracy of 100%, but this could be due in part to the memoization effect between experiments, this is particularly noticeable on the k-fold experiments.

The reason we expected the best results for the CNN is that we took measures to reduce overfitting, first, we used 2 dropout layers to reduce overfitting, by regularizing the model we try and encourage each neuron to be more independently useful. Using the Keras Dataflow class, we produced an augmented dataset, this allowed us to increase the representation of the least represented sign classes, by providing random variance within a zoom, rotation and shift range, effectively providing an infinite stream of data to train with.

For further experiments, it would be interesting to investigate the effects of changing the parameters of the ImageDataGenerator class to train with a wider variety of data.

From Figure 8, we can see that using an augmented dataset reduces overfitting and delays the convergence epoch to later in the training process.

augmented	epochs	Label_class	accuracy	best_epoch	best_val_lo	epoch	loss	num_classes	val_accuracy	val_loss
true	50	beware cycle route	0.9998	21	0.001906	49	0.001184	2	0.9945	0.2462
true	50	beware children	0.9993	23	0.004586	49	0.007824	2	0.9961	0.1432
true	50	beware pedestrian	0.9999	12	2.089	49	0.0002245	2	0.9903	98.18
true	50	right turn	0.9997	44	4.707e-9	49	0.001879	2	0.9987	0.003801
true	50	left turn	1	40	3.858e-11	49	1.897e-8	2	1	0
true	50	speed limit 70	0.9995	9	0.008729	49	0.005162	2	0.9958	0.04215
true	50	speed limit 60	0.9998	33	0.001105	49	0.001557	2	0.9981	0.005855
true	50	speed limit 50	0.9994	5	0.01733	49	0.003941	2	0.9906	0.06074
true	50	speed limit 30	0.9985	40	0.003547	49	0.00838	2	0.9951	0.03181
true	50	speed limit 20	0.9996	20	5.263e-7	49	0.003386	2	1	0.0000016
true	50	All Classes	0.995	28	0.08317	49	0.0244	10	0.9505	0.7379
false	50	beware cycle route	1	49	0.0000404	49	2.069e-8	2	1	0.0000404
false	50	beware children	1	5	0.009945	49	0	2	0.9974	0.03543
false	50	beware pedestrian	1	5	0.7574	49	1.230e-11	2	0.9903	1.292
false	50	right turn	1	7	3.086e-10	49	0	2	0.9997	0.0003725
false	50	left turn	1	6	3.858e-11	49	0	2	1	0
false	50	speed limit 70	1	7	0.05147	49	0	2	0.9958	0.1444
false	50	speed limit 60	1	44	0.003328	49	0	2	0.999	0.02362
false	50	speed limit 50	0.9999	4	0.059	49	0.003891	2	0.9926	0.2552
false	50	speed limit 30	0.9997	26	0.03957	49	0.002865	2	0.9964	0.07073
false	50	speed limit 20	1	17	1.157e-10	49	0	2	0.9984	0.02727
false	50	All Classes	0.9989	2	0.3193	49	0.01266	10	0.979	1.332

Figure 8: CNN results with baseline testing data, accuracy is training accuracy, and val_accuracy is testing accuracy.

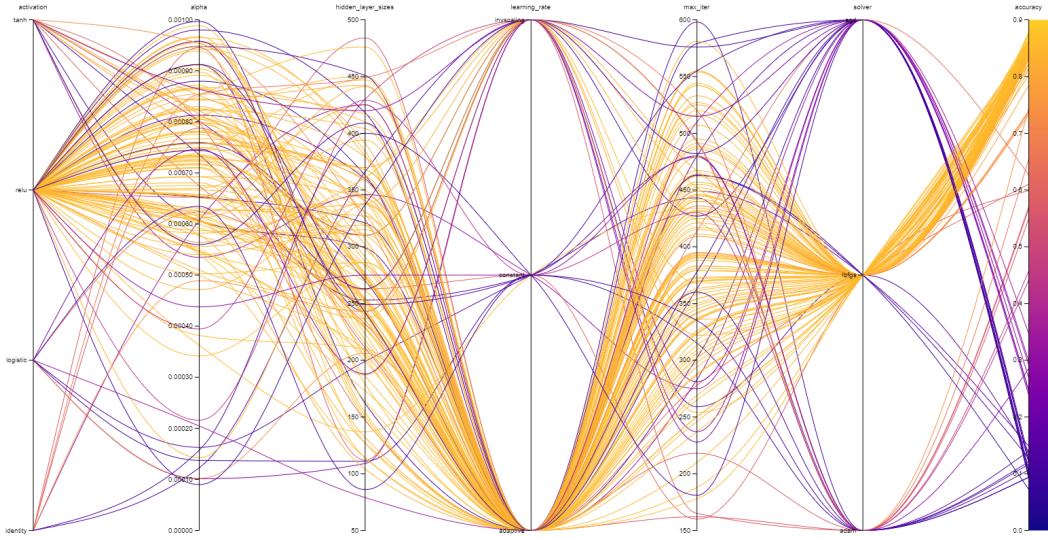
4 Methodology for Optimising hyperparameters

Our hyperparameter optimization used Bayesian Optimization [3]. This gaussian process models the function and then chooses parameters to optimize by determining the probability of improvement according to the accuracy metric.

We opted to use the ‘all classes’ dataset for these sweeps because in our preliminary experiments we noticed the worst prediction accuracy with 10 classes.

Figure 9 shows a high-level overview of the MLP sweep, dark blue lines indicate lower accuracy and yellow had the highest accuracy. (See Section 6.2)

Figure 9: Multilayer Perceptron



For the evaluation of our hyperparameters, we gathered data on the linear correlation between each hyperparameter and their accuracy. The feature importance is calculated from training a random forest with hyperparameters and its metrics. [4]

The feature importance metric is useful to identify the features that have a larger impact on the output of the model, coupled with correlation we can find the direction of the relationship. It may be the case that a very important feature is extremely bad for the result, in this case, we will see a strong negative correlation, this information helps us hone in on the best hyperparameters to use to maximise accuracy.

Using the accuracy metric however, caused the models to instead maximise the hyperparameters tendency to overfit this very imbalanced dataset.

5 Variation in performance with varying learning parameters in Decision Trees

5.1 J48

Discussed below is the hyperparameter [5] importance concerning the performance of J48 trees:

max_depth - The maximum depth of the tree.

min_impurity_decrease - Splits a node if this split induces a decrease of the impurity greater than or equal to this value.

min_samples_leaf - the minimum number of samples to be considered a leaf.

min_weight_fraction_leaf - The minimum number of samples required to be at a leaf node.

criterion_value -['Gini', 'Entropy']: Measures the quality of the split.

max_features - ['auto', 'log2', 'sqrt'] : max features considers the number of features when looking for the best split. Auto just picks the best result so had the same result as log2.

splitter - ['best', 'random']: The strategy used to choose the split at each node.

Parameter	Type	Conclusion
max_depth	int	Provided low importance value of 0.012 and provided some correlation 0.084.
min_impurity_decrease	float	Small importance value of 0.051 and a strong negative correlation of -0.241.
min_samples_leaf	int or float	Low importance of 0.014 and provided a tiny positive correlation of 0.007.
min_weight_fraction_leaf	float	Gives a strong negative correlation in terms of accuracy, meaning the higher the min_weight_fraction_leaf value, the lower the accuracy.
criterion	string	Both parameters provided a very low importance value of 0.003. Entropy had a tiny negative correlation of -0.022 and Gini a positive correlation of 0.022.
max_features	string	Log2 provided an importance value of 0.133 and a high negative correlation of -0.372. Log2 provided tiny importance of 0.003 but a relatively good correlation of 0.196.
splitter	string	'best' had an importance of 0.048 and random 0.044. 'best' has a positive correlation of 0.051 whereas random -0.051.

See Parameter Importance (Figure D.1)

5.2 Random Forest

Discussed below is the hyperparameter [6] importance concerning the performance of Random Forest trees:

min_samples_split - The minimum number of samples required to split an internal node

min_samples_leaf - The minimum number of samples required to be at a leaf node.

n_estimators - The number of trees in the forest.

min_weight_fraction_leaf - The minimum number of samples required to be at a leaf node.

max_features - The number of features to consider when looking for the best

criterion - The function to measure the quality of a split. Supported criteria are “gini” for the Gini impurity and “entropy” for the information gain.

Parameter	Type	Conclusion
max_features	string	Contains the options ”auto”, ”sqrt” and ”log2”. We discovered that ”sqrt” has a higher accuracy overall, the accuracy of ”log2” varies between the lower end and the median accuracy value.
min_samples_split	int or float	The minimum samples required to split a node has very little impact on accuracy.
criterion	string	Gives a perfect negative correlation with respect to accuracy. Correlation values being [Gini = -0.404], [Entropy = 0.404].
n_estimators	int	This is defined as the number of trees in the forest, it seems to have very little correlation but high importance.
min_samples_leaf	int or float	Gives a strong negative correlation in terms of accuracy, meaning the higher minimum samples at a leaf node, the lower the accuracy.
min_weight_fraction_leaf	float	Has a somewhat positive correlation to accuracy. e.g. total weight required at a leaf node varies between 76% and 89% accuracy

See Parameter Importance (Figure E.1)

6 Variation in performance with varying learning parameters in Neural Networks

6.1 Linear Classifier

For a linear classifier, we have decided to use logistic regression. [7] Shown below is the hyperparameter importance concerning the performance of Logistic Regression.

The accuracy fluctuates depending on which hyperparameters are used. Logistic Regression has a ‘solver’ hyperparameter in sklearn which is the algorithm to use in the optimization. We have tested the following: ‘newton-cg’, ‘lbfgs’, ‘liblinear’, ‘sag’, ‘saga’. Sag and Saga give the best results with 89% accuracy. Liblinear isn’t far behind with 88%. Other solvers have a negative impact on the performance, this can be seen in the Solver parameter importance table where their correlation values are negative.

Additionally ‘sag’ and ‘saga’ guarantee fast convergence on features with approximately the same scale [R41] which is the case with our datasets as pixel values have the same scale. This can be seen as the number of iterations between 200-800 don’t change much in the accuracy and from the scatter chart of the accuracy of the Logistic Regression vs the number of sweeps created.

Parameter	Type	Conclusion
C	float	The smaller this value is the stronger the regularization. For this we settled on 0.08027. Although this parameter did not make a big difference and provided similar results for values between 0.8 and 1.2
fit_intercept	bool	From testing we have found out that for this dataset adding a bias works better than not having bias thus True
max_iter	int	We settled on 342 maximum number of iterations taken for the solver to converge. This parameter does not have a big importance but we found 342 gives accurate results.
solver	string	Contains the options ‘newton-cg’, ‘lbfgs’, ‘liblinear’, ‘sag’, ‘saga’, default=’lbfgs’. ‘Saga’ gives the best accuracy overall. Solver is the most influential hyperparameter as we can see from the importance and correlation table
tol	float	tolerance for stopping criteria which tells the algorithm to stop searching when some tolerance is achieved. This parameter did not make a difference, we settled on 0.0002277

See Parameter Importance (Figure F.1)

6.2 Multilayer Perceptron

Our sweep on MLP [8] indicated that the relu activation function provided the highest accuracy metric for all classes, our search computed this as the most important metric. However, using this we observed severe overfitting on the speed limit signs in all experiments.

activation - Activation function for the hidden layer.

alpha - L2 penalty (regularization term) parameter.

hidden_layer_sizes - Number of neurons in each hidden layer.

max_iter - Maximum number of iterations. The solver iterates until convergence

solver - The solver for weight optimization.

learning_rate - The learning rate controls how quickly the model is adapted to the problem.

Parameter	Type	Conclusion
alpha	float	When alpha increases the loss increases therefore we went with a relatively low number of 0.0005508.
solver	string	'lbfgs' solver did provide the best accuracy consistently when compared to other solvers such as 'adam' and 'sgd' which supports the statement from [8] staying that 'Lbfgs' works best for small datasets.
max_iter	int	Maximum number of iterations. We used 520 as it gives good results but that is the max so it converges earlier, as long as a value above 200 is used will give high accuracy.
activation	string	Contains the options: 'identity', 'logistic', 'tanh', 'relu'. We found relu worked best for our use case
learning_rate	string	'adaptive' achieves the highest accuracy while, 'constant' and 'invscaling' vary widely.
hidden_layer_sizes	tuple	337 layers worked well for us. But this parameter is not very important as long at its above a fairly high value such as 200.

See Parameter Importance (Figure G.1)

6.3 Further experiments

We noticed extreme overfitting when working with most classifiers after running our hyperparameter search, next we decided to run a new search, using one of our worst performers, the Multi-Layer perceptron because we felt it had the most room for improvement. Therefore we decided to experiment with it further by running a more complex sweep. To do this we changed the optimisation metric to minimize loss and randomly varied the classification labels throughout the search, our goal was to see if we could find a single set of parameters that resulted in a better recall, precision, f1, accuracy and loss.

6.3.1 Turn left label

Looking closer at the turn left label results a few patterns emerge, the hidden layer size was awarded the highest importance with a low positive correlation to a lower loss, and the learning rate hyperparameter set to constant had the highest overall feature importance with a negative correlation to loss of -0.301.

6.3.2 Beware Pedestrian Crossing - Looking closer

We found a few instances in the sweep providing a Precision and Recall Score of 1 10. The first observation we made concerned the solver hyperparameter. The majority of these results occurred with the ‘lbfgs’ solver, however, ‘sgd’ was able to achieve these recall and precision metrics also but with a much higher loss between 0.03 and 1.4, whereas ‘lbfgs’ consistently had a loss of less than 0.0016. The variance in loss when the ‘identity’ activation function was used was extreme, from 0.00009 to 2.28e+22, this indicates this activation would be an extremely unreliable choice.

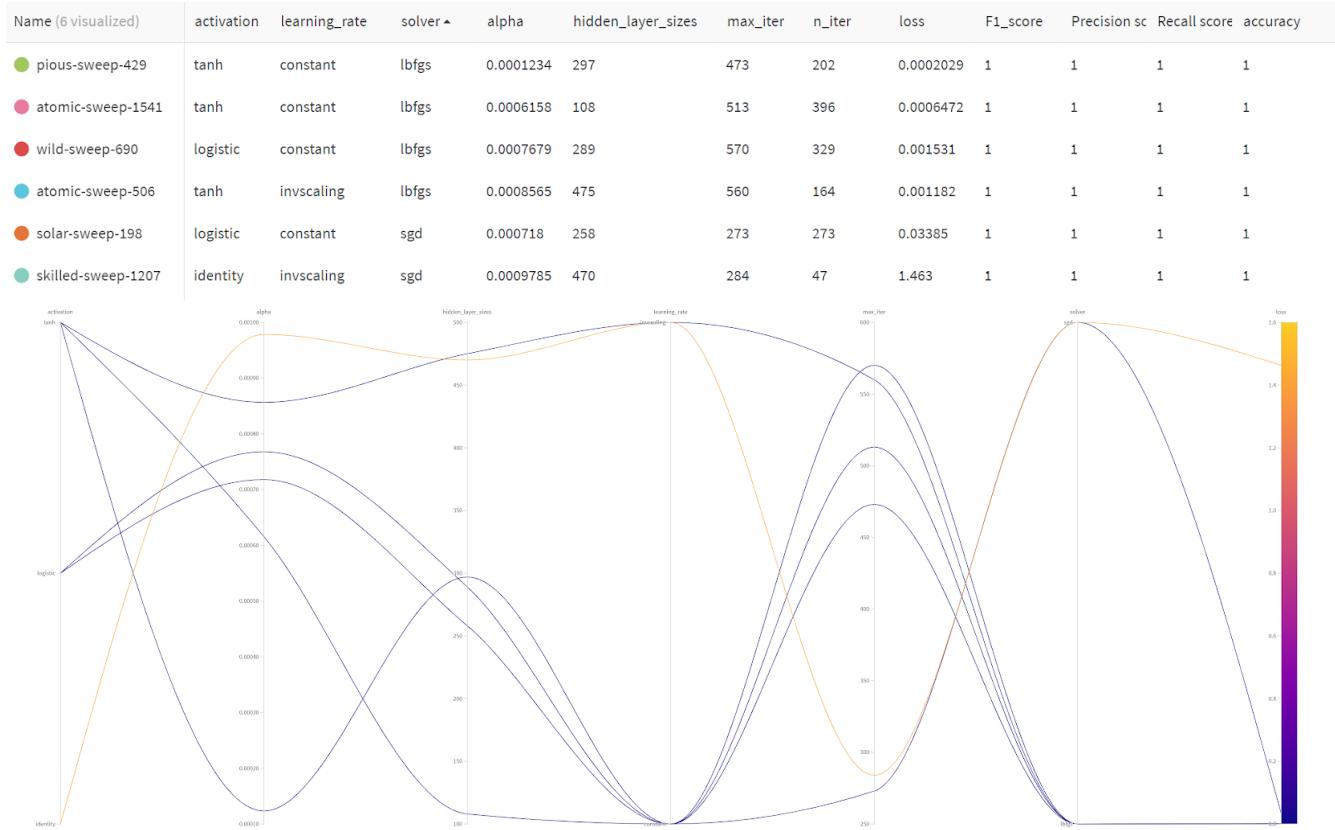


Figure 10: MLP Pedestrian Crossing, table of hyperparameter sweep results with corresponding line chart

During the experiment, we noticed the worst predictions with 0 precision and recall (See Fig 11) that MLP results almost always use either ‘adam’ or ‘sgd’ solver and tend to iterate around 20 times before terminating in most cases. We suspect this due to the algorithm being stuck in local maxima and the convergence conditions are satisfied, thus the algorithm terminates early. We observed that the majority of class labels where this occurs are speed signs.

Name (50 visualized)	activation	hidden_lay	learning_rate	max_iter	n_iter	solver	alpha	accuracy	loss
● ► Class label: "speed limit 30" 6	["relu","identity","tanh","logistic"]	288.5	["invscaling","constant"]	403.167	15.5	["adam","sgd"]	0.0005979	0.5048	2.1382713
● ► Class label: "speed limit 20" 6	["identity","tanh","logistic"]	261.667	["constant","adaptive"]	314.5	21.833	["sgd","adam"]	0.0005876	0.5024	1.2573099
● ► Class label: "beware pedestrian crossing" 2	["tanh","identity"]	111	["invscaling","constant"]	412	22.5	["adam","sgd"]	0.0004747	0.4843	401085101
● ► Class label: "speed limit 50" 7	["identity","tanh","logistic"]	288.143	["constant","invscaling","adaptive"]	406.571	15.286	["sgd","adam"]	0.0006154	0.5138	9.1698027
● ► Class label: "speed limit 70" 7	["logistic","tanh","identity"]	388	["invscaling","adaptive","constant"]	352.286	19	["adam","sgd"]	0.0005835	0.5119	5.6478075
● ► Class label: "speed limit 60" 6	["relu","identity"]	311	["invscaling","constant"]	251.333	23.833	["adam","sgd"]	0.000272	0.5209	3.5168506
● ► Class label: "left turn" 3	["identity","tanh"]	386.333	["constant","invscaling"]	288.667	22.667	["sgd","adam"]	0.0006752	0.4796	94029864:
● ► Class label: "beware children" 6	["tanh","identity"]	224.167	["constant","adaptive","invscaling"]	359.167	17.167	["adam","lbfgs","sgd"]	0.0004487	0.4846	1.0665595
● ► Class label: "beware cycle route ahead" 2	["tanh","relu"]	293	invscaling	376	27	["adam","lbfgs"]	0.0001605	0.4804	9.187
● ► Class label: "right turn" 5	identity	176.8	constant	352.4	12.2	sgd	0.0003253	0.5084	3.3483492

Figure 11: MLP sweep results - grouping by class label where precision & recall are 0.

6.3.3 Conclusions about the extended MLP sweep

Overall, the experiment made it clear that each test set would benefit from individual hyperparameter profiles customised to the data, the only notable correlation with a minimized loss we observed was the 'lbfgs' solver, with a positive correlation for recall and precision scores of 0.317, otherwise all metrics demonstrated correlations between -0.2 and 0.078 12.

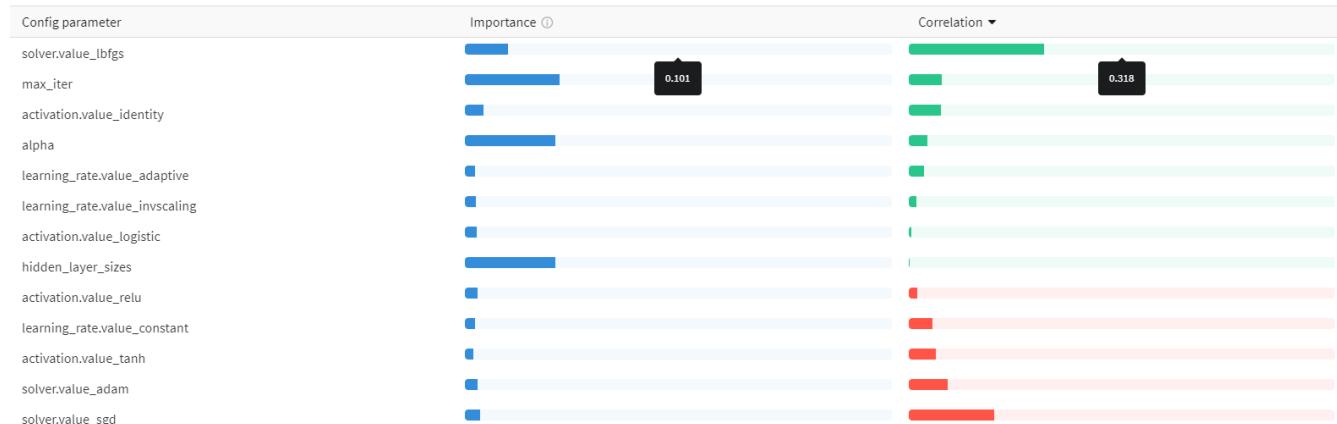


Figure 12: MLP - Parameter importance with respect to Recall score

7 Variation in performance according to different metrics

7.1 J48 Experiments

J48 with default test set cycle route ahead binary classification gave a very high precision (0.95) and a really low recall (0.2) Precision shows the ratio of true positives to all the positives, therefore in our use case we have a ratio of signs that the classifier classified correctly to the total the classifier thought are positive.

Using cycle route ahead dataset as an example, for all signs that are actually cycle route ahead signs the recall tells us how many the classifier classified as cycle route ahead signs. Therefore J48 is very picky and doesn't think many signs are cycle route ahead signs. Pretty much all signs it thinks are cycle route ahead signs are indeed cycle route ahead signs. However, it also misses a lot of cycle route ahead signs, because it is very picky.

With the default test set, J48 provides a precision, recall and f1 scores of 0 for the right turn, left turn, beware pedestrian crossing and speed limit 20 classes the reason is due to 0 true positive classifications. When we investigate why this is for the beware pedestrian crossing class, we noticed the data has a heavy skew of class proportions towards 'No' with very few instances of 'Yes' found. In fact, this is the case for all of the class labels with scores of 0 for precision, recall and f1.

7.2 Random Forest & MLP

Both random forest and MLP suffered overfit issues particularly with binary class labels; they scored 0 for precision, recall, and F1 in all experiments. We hypothesize this is caused by the same reasons described in Section 2. Likewise, all classes perform poorly the model classifies every class as exclusively either speed limit 30, 50, and 70 in the training and testing data sets, the k-fold experiments simply classifies all classes as speed limit 50; the most well-represented class in our dataset.

7.3 Logistic Regression

Logistic Regression is excellent at predicting classes for the dataset, which contains all signs, which is not easy to do. Logistic Regression has relatively good precision and recall when compared to other classifiers with 89 for both. This means it is very good at predicting all different signs (See Figure 13).

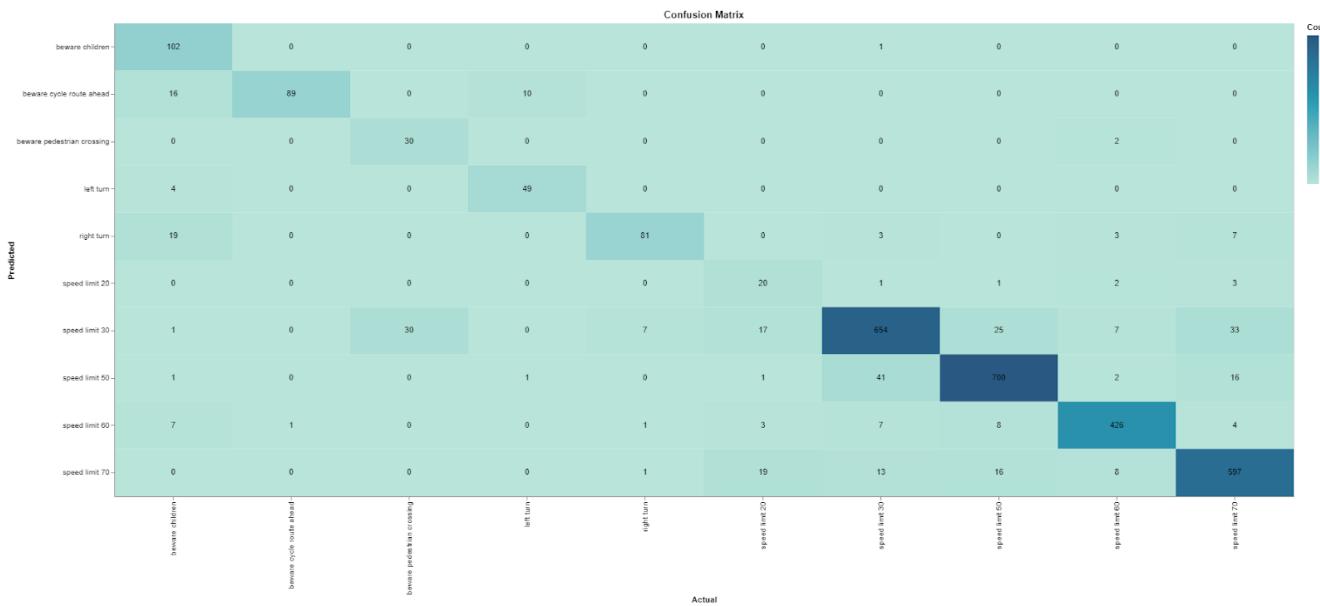


Figure 13: Logistic regression confusion matrix all classes

Appendices

A Workload

A.1 Workload split

Team member	Involvement
Lewis Wilson	Contributed to the experimenting and analysis of J48, Report writing.
Chun Man	Contributed to the experimenting and analysis of Random Forest, Report writing.
Sam Fay-Hunt	Contributed to the report, coded all the automated experiments, weights and biases logging and integration, the interface used by each sklearn model, and the entire CNN.
Kamil Szymczak	Contributed to the experimenting and analysis of Logistic Regression, Report writing.

As a team, we are happy with everyone's contributions to the project. All team members were punctual and showed up to all scheduled meetings. Sam took the lead as project manager throughout the project delegating the workload and providing support to others.

B References

References

- [1] (). “Weights & Biases,” W&B, [Online]. Available: <https://wandb.ai/lavanyashukla/vega-plots/reports/Log-ROC-PR-curves-and-Confusion-Matrices-with-W-B--Vmlldzo3NzQ3MQ> (visited on 12/08/2020).
- [2] Y. Zhao and Y. Zhang, “Comparison of decision tree methods for finding active objects,” *Advances in Space Research*, vol. 41, no. 12, pp. 1955–1959, Jan. 2008, ISSN: 02731177. DOI: 10.1016/j.asr.2007.07.020. arXiv: 0708.4274. [Online]. Available: <http://arxiv.org/abs/0708.4274> (visited on 12/07/2020).
- [3] (). “Configuration,” [Online]. Available: <https://docs.wandb.com/sweeps/configuration> (visited on 12/07/2020).
- [4] (). “Parameter Importance,” [Online]. Available: <https://docs.wandb.com/app/features/panels/parameter-importance> (visited on 12/07/2020).
- [5] (). “Sklearn.tree.DecisionTreeClassifier — scikit-learn 0.23.2 documentation,” [Online]. Available: <https://scikit-learn.org/stable/modules/generated/sklearn.tree.DecisionTreeClassifier.html#sklearn.tree.DecisionTreeClassifier> (visited on 12/07/2020).
- [6] (). “3.2.4.3.1. sklearn.ensemble.RandomForestClassifier — scikit-learn 0.23.2 documentation,” [Online]. Available: <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestClassifier.html> (visited on 12/07/2020).
- [7] (). “Sklearn.linear_model.LogisticRegression — scikit-learn 0.23.2 documentation,” [Online]. Available: https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html (visited on 12/07/2020).
- [8] (). “Sklearn.neural_network.MLPClassifier — scikit-learn 0.23.2 documentation,” [Online]. Available: https://scikit-learn.org/stable/modules/generated/sklearn.neural_network.MLPClassifier.html (visited on 12/07/2020).

C Links to experimental results

The following is a list of the links to all of our experimental data.

C.1 J48

- Kfold
- Baseline test
- 4k instances moved to test
- 9k instances moved to test
- Hyperparameter sweep

C.2 Random Forest

- Kfold
- Baseline test
- 4k instances moved to test
- 9k instances moved to test
- Hyperparameter sweep

C.3 MLP

- Kfold
- Baseline test
- 4k instances moved to test
- 9k instances moved to test
- Initial hyperparameter sweep
- Enhanced hyperparameter sweep

C.4 LR

- Kfold
- Baseline test
- 4k instances moved to test
- 9k instances moved to test
- Hyperparameter sweep

C.5 CNN

- Kfold
- Baseline test
- 4k instances moved to test
- 9k instances moved to test

D J48

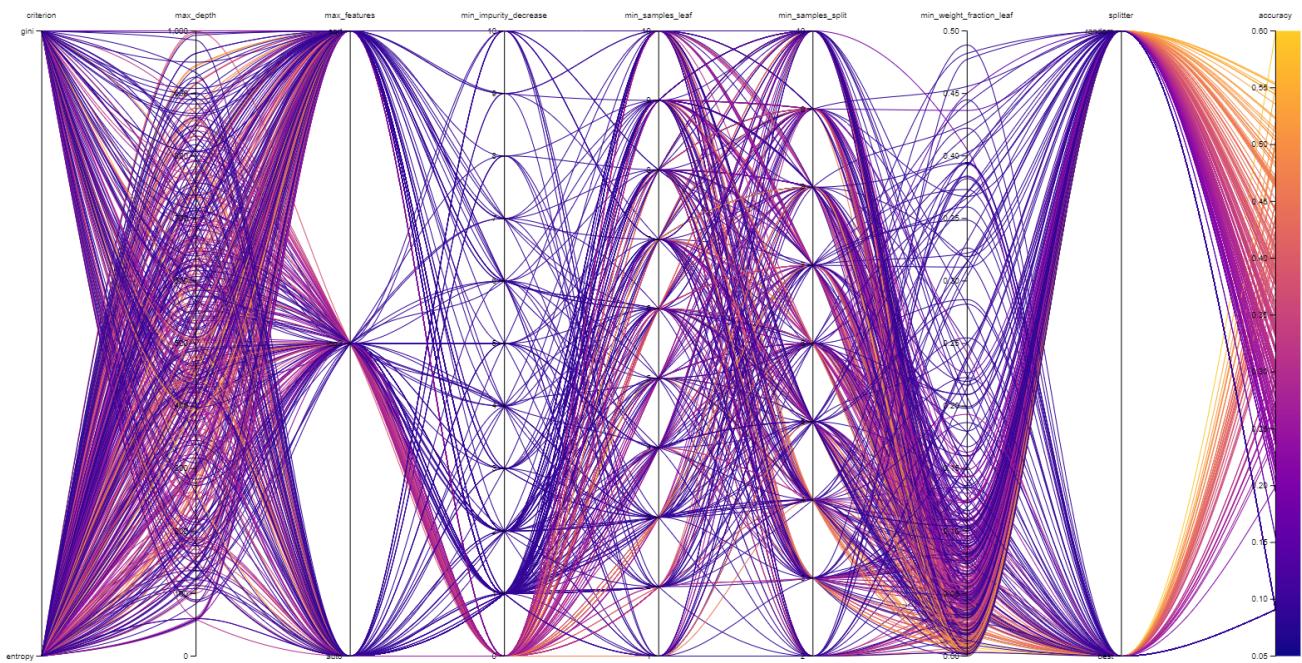
D.1 J48 Parameter Importance

Table 2: J48 Parameter metrics

Parameter Config	Importance	Correlation
min_weight_fraction_leaf	0.668	-0.666
min_impurity_decrease	0.052	-0.241
min_samples_leaf	0.016	0.007
max_depth	0.012	0.084
min_samples_split	0.009	-0.108

Parameter Config	Importance	Correlation
max_features.value_log2	0.133	-0.372
max_features.value_sqrt	0.003	0.196
splitter.value_best	0.049	0.051
splitter.value_random	0.043	-0.051
criterion.value_entropy	0.002	-0.022
criterion.value_gini	0.002	0.022

Figure 14: J48 Parameters



E Random Forest

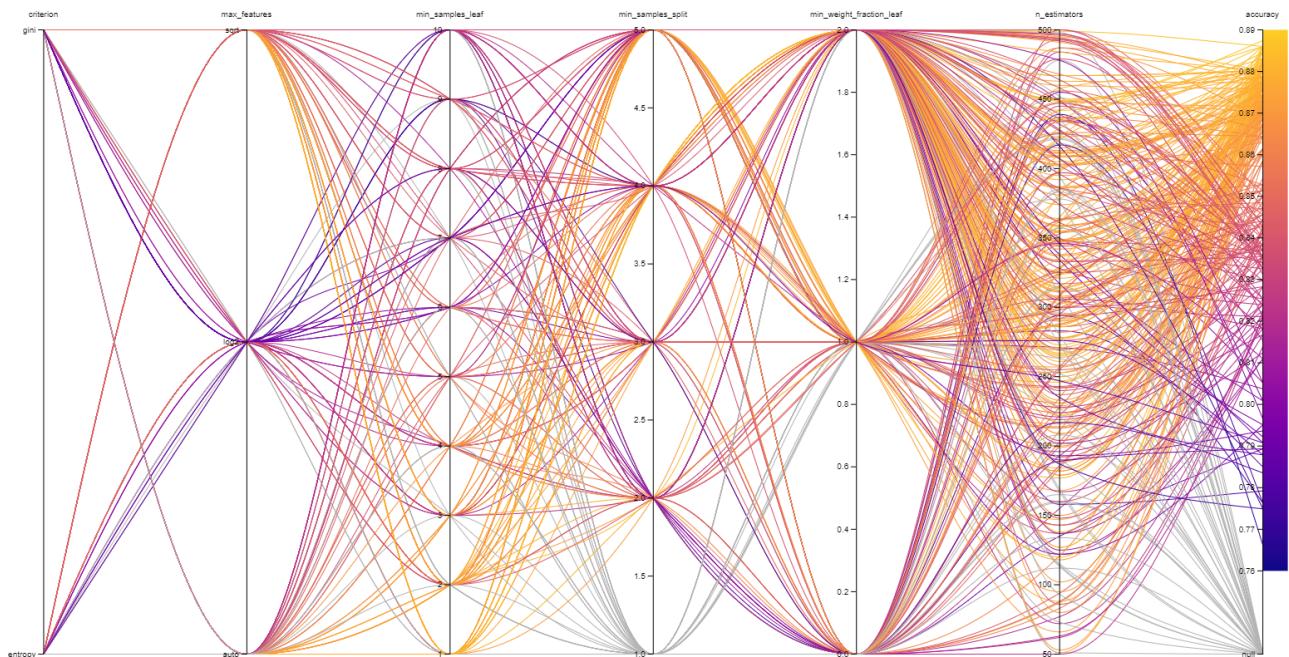
E.1 Random Forest Parameter Importance

Table 3: Random Forest Parameter metrics

Parameter Config	Importance	Correlation
min_samples_split	0.005	0.306
min_samples_leaf	0.375	-0.725
n_estimators	0.016	0.092
min_weight_fraction_leaf	0.013	0.123

Parameter Config	Importance	Correlation
max_features.value_sqrt	0.001	0.563
max_features.value_log2	0.565	-0.752
criterion.value_entropy	0.012	0.404
criterion.value_gini	0.012	-0.404

Figure 15: Random Forest Parameters



F Linear Classifier

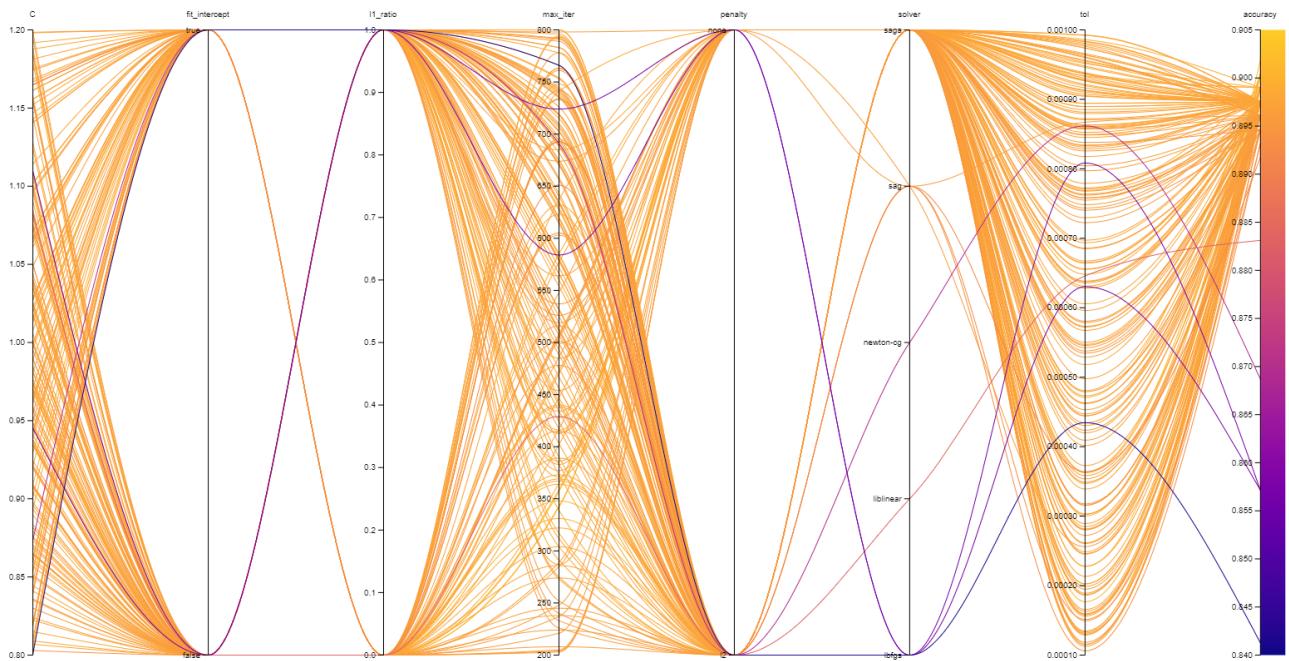
F.1 Linear Classifier Parameter Importance

Table 4: Linear Classifier Parameter metrics

Solver Parameter Config	Importance	Correlation
lbfgs	0.786	-0.869
newton-cg	0.159	-0.271
liblinear	0.042	-0.037
saga	0.009	0.691
sag	0.003	0.170

Config Parameter Config	Importance	Correlation
max-iter	0.194	-0.634
l1-ratio	0.124	-0.510
fit-intercept	0.058	0.367
tol	0.046	-0.174
C	0.031	0.260

Figure 16: Random Forest Parameters



G Multilayer Perceptron

G.1 Multilayer Perceptron Parameter Importance

Table 5: MLP Parameter metrics

Parameter Config	Importance	Correlation
hidden_layer_sizes	0.095	-0.101
max_iter	0.091	0.072
alpha	0.061	0.171

Parameter Config	Importance	Correlation
solver.value_lbfsgs	0.530	0.728
solver.value_adam	0.027	-0.245
solver.value_sgd	0.024	-0.640
activation.value_identity	0.098	-0.169
activation.value_relu	0.033	0.301
activation.value_tanh	0.018	-0.035
activation.value_logistic	0.006	-0.270
learning_rate.value_adaptive	0.012	0.507
learning_rate.value_constant	0.004	-0.442
learning_rate.value_invscaling	0.001	-0.224

Figure 17: Multilayer Perceptron Parameters

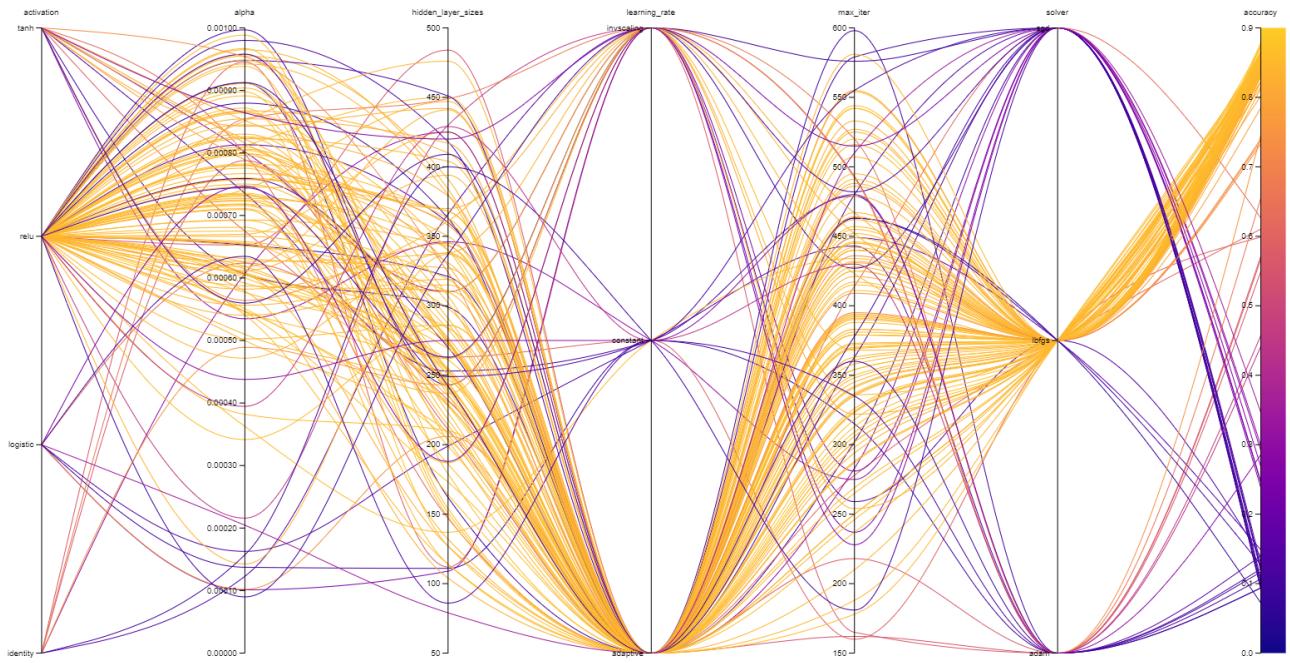


Figure 18: Data with evenly split groups

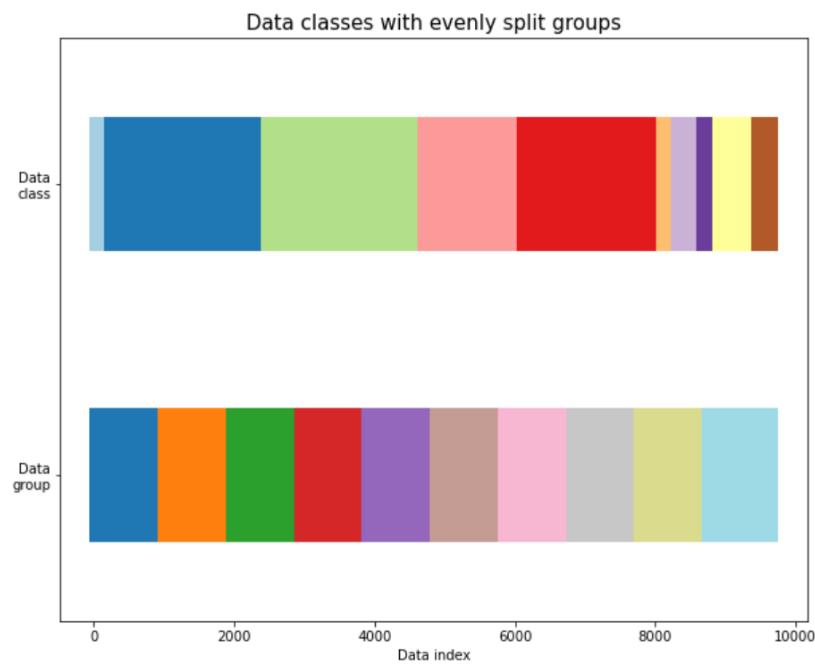


Figure 19: KFold

