

# Data Mining & Machine Learning F20DL

Group 4

Lewis Wilson, Sam Fay-Hunt, Kamil Szymczak, Chun Man

December 6, 2020

# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Variation in performance with size of the training and testing sets</b>	<b>2</b>
2.1	Methodology . . . . .	2
2.2	Turn left deep dive . . . . .	2
2.3	Original train & test sets . . . . .	3
2.4	4000 moved . . . . .	3
2.5	9000 moved . . . . .	3
2.6	J48 Tree Classifier . . . . .	3
<b>3</b>	<b>Variation in performance with the change in the learning paradigm (Decision Trees versus Neural Nets)</b>	<b>4</b>
<b>4</b>	<b>Variation in performance with varying learning parameters in Decision Trees</b>	<b>5</b>
4.1	Optimising Hyperparameters . . . . .	5
4.2	J48 . . . . .	6
4.3	Random Forest . . . . .	7
<b>5</b>	<b>Variation in performance with varying learning parameters in Neural Networks</b>	<b>8</b>
5.1	Linear Classifier . . . . .	8
5.2	Multilayer Perceptron . . . . .	9
5.3	Convolutional Neural Networks . . . . .	10
5.3.1	Critical limitation of CNN experimental methodology: . . . . .	10
<b>6</b>	<b>Variation in performance according to different metrics (TP Rate, FP Rate, Precision, Recall, F Measure, ROC Area)</b>	<b>11</b>
<b>Appendices</b>		<b>12</b>
<b>A</b>	<b>Appendix A</b>	<b>12</b>
A.1	Workload split . . . . .	12

<b>B J48</b>	<b>13</b>
B.1 J48 Parameter Importance . . . . .	13
<b>C Random Forest</b>	<b>14</b>
C.1 Random Forest Parameter Importance . . . . .	14
<b>D Linear Classifier</b>	<b>15</b>
D.1 Linear Classifier Parameter Importance . . . . .	15
<b>E Multilayer Perceptron</b>	<b>16</b>
E.1 Multilayer Perceptron Parameter Importance . . . . .	16

# 1 Introduction

For both decision trees and neural networks, we used Weights and Biases (<https://wandb.ai/home>) to visualise these experiments which we used to record and organise the experiment data shown the graphs shown in this document.

## 2 Variation in performance with size of the training and testing sets

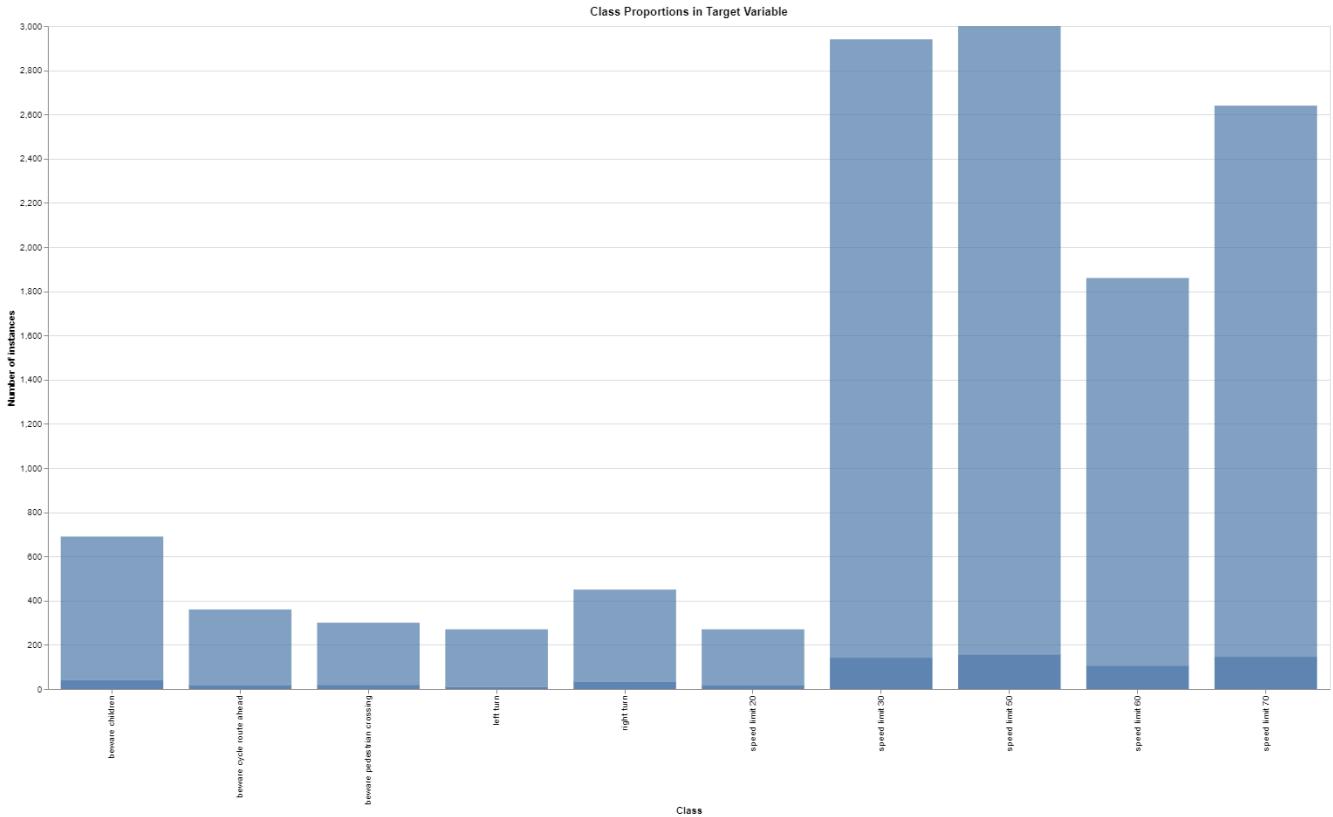
### 2.1 Methodology

We used the same seed when moving the data from the training set to the testing set, so all experiments for a given classifier used the same training and testing data.

### 2.2 Turn left deep dive

The turn left dataset has an extremely small number of instances in its training dataset for 9000 moved to testing. Thus we would expect the worst prediction accuracy, in this case, the situation was so bad that our accuracy scores became misleading. Because the representation of ‘turn left’ was so poor in our dataset most of the classifiers simply predicted that there were no instances of turn left in the dataset and recorded a deceiving 97% accuracy (because 97% of the dataset is not in fact turn left). In fact this problem appeared for all binary classifications with the Random Forest classifier. The only classifiers that did not exhibit this behaviour were the Logistic Regression (LR) and Convolutional Neural Network classifiers (CNN). CNN scored a staggering 99.9% accuracy with the same minuscule training data, while LR managed 98.3% with 200 false negatives, 5 false positives and 60 true positives.

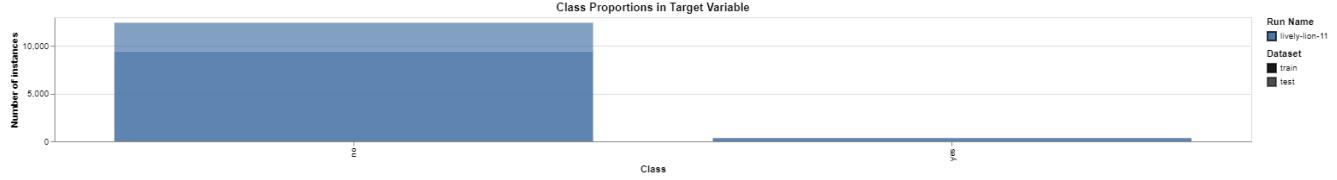
Figure 1: Class proportions in target variable



An example distribution of labels between train (dark blue) and test (light blue), for the Beware cycle route ahead class.

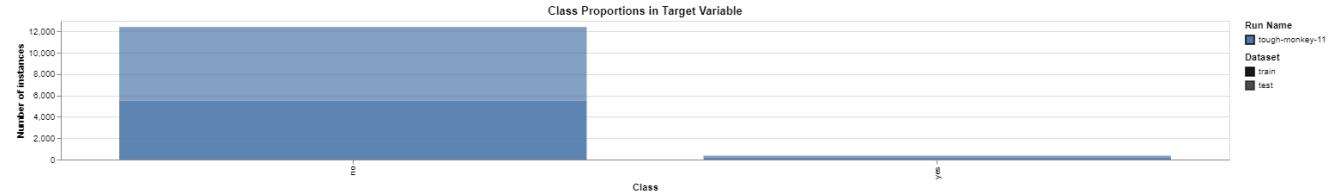
## 2.3 Original train & test sets

Figure 2: Class proportions in target variable



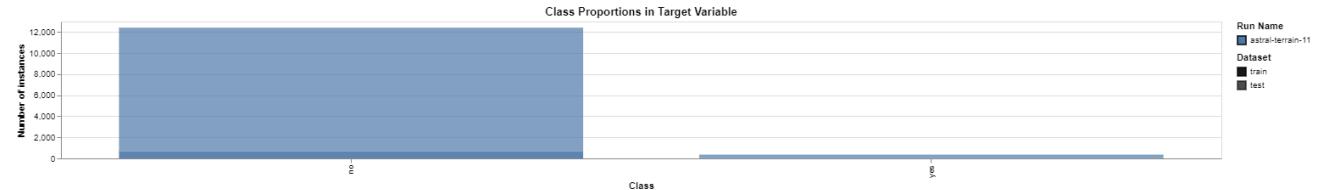
## 2.4 4000 moved

Figure 3: Class proportions in target variable



## 2.5 9000 moved

Figure 4: Class proportions in target variable



## 2.6 J48 Tree Classifier

**As expected, moving randomly only 4000 instances from original training to testing set gives better accuracy when compared to moving 9000 instances as the classifier has less data to learn from and sees more new never seen data. The difference was quite big varying from 1% to as much as 9%, the only exception was the speed limit 20 sign that achieved the same accuracy whereas the opposite can be said about 30 and 50 speed limit signs that had a huge difference of 9% decrease in accuracy when compared 4000 moved to 9000 moved. Despite different accuracies the confusion matrix remained the same having similar proportions.**

### **3 Variation in performance with the change in the learning paradigm (Decision Trees versus Neural Nets)**

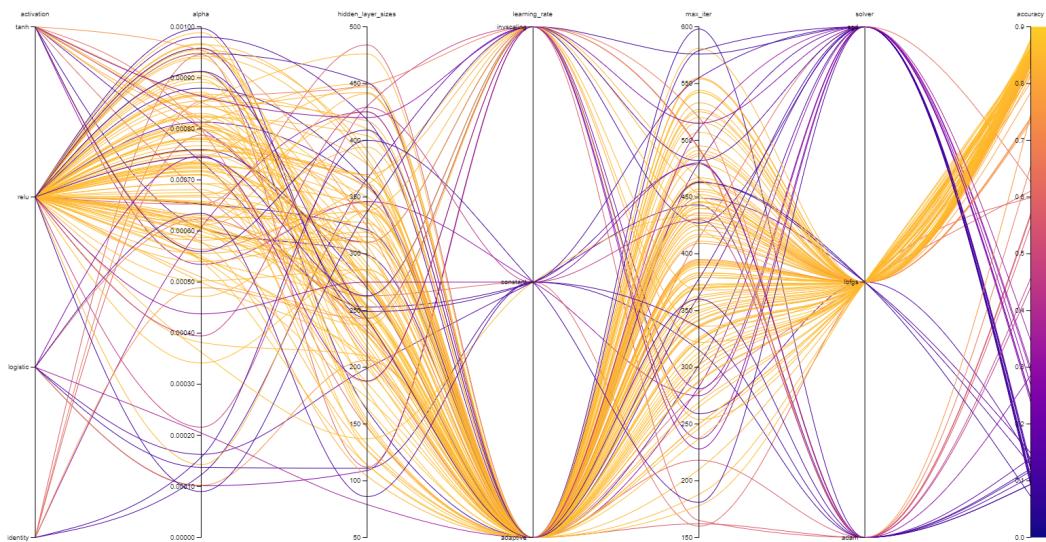
- Decision trees performed comparatively poorly, particularly when the training data was reduced.

## 4 Variation in performance with varying learning parameters in Decision Trees

### 4.1 Optimising Hyperparameters

We used a sweeping tool to try a large number of hyperparameters with a bayesian search strategy that uses a Gaussian process to model the function and then chooses parameters to optimize the probability of improvement according to the accuracy of the test set. **NEEDS REPHRASE!** <https://docs.wandb.com/sweeps/configuration#search-strategy>. We opted to use the ‘all classes’ dataset for these sweeps because in our preliminary experiments we noticed the worst prediction accuracy with 10 classes.

Figure shows high level overview of the MLP sweep, dark blue lines indicate lower accuracy and yellow had the highest accuracy. (See 5.2)



To evaluate the hyperparameters we collected and evaluated the linear correlation between the hyperparameter and the accuracy, and also a feature importance metric provided by weights and biases, this metric is computed by training a random forest with the hyperparameters as inputs and the accuracy as the target output, the feature importances are then extracted from the random forest. <https://docs.wandb.com/app/features/panels/parameter-importance>

The feature importance metric is useful to identify the features that have a larger impact on the output of the model, coupled with correlation we can find the direction of the relationship. It may be the case that a very important feature is extremely bad for the result, in this case we will see a strong negative correlation, this information helps us hone in on the best hyperparameters to use.

## 4.2 J48

**max\_depth** - The maximum depth of the tree.

**min\_impurity\_decrease** - Splits a node if this split induces a decrease of the impurity greater than or equal to this value.

**min\_samples\_leaf** - the minimum number of samples to be considered a leaf.

**min\_weight\_fraction\_leaf** - The minimum number of samples required to be at a leaf node.

**criterion\_value** -['Gini','Entropy']: Measures the quality of the split. **max\_features** - ['auto','log2','sqrt'] : max features considers the number of features when looking for the best split. Auto just picks the best result so had the same result as log2.

**Splitter** - ['best', 'random']: The strategy used to choose the split at each node.

Parameter	Type	Conclusion
max_depth	int	Provided low importance value of 0.012 and provided some correlation 0.084.
min_impurity_decrease	float	Small importance value of 0.051 and a strong negative correlation of -0.241.
min_samples_leaf	int or float	Low importance of 0.014 and provided a tiny positive correlation of 0.007.
min_weight_fraction_leaf	float	Gives a strong negative correlation in terms of accuracy, meaning the higher the min_weight_fraction_leaf value, the lower the accuracy.
criterion	string	Both parameters provided a very low importance value of 0.003. Entropy had a tiny negative correlation of -0.022 and Gini a positive correlation of 0.022.
max_features	string	Log2 provided an importance value of 0.133 and a high negative correlation of -0.372. Log2 provided tiny importance of 0.003 but a relatively good correlation of 0.196.
splitter	string	'best' had an importance of 0.048 and random 0.044. 'best' has a positive correlation of 0.051 whereas random -0.051.

See Parameter Importance (Figure B.1)

### 4.3 Random Forest

**min\_samples\_split** - The minimum number of samples required to split an internal node

**min\_samples\_leaf** - The minimum number of samples required to be at a leaf node.

**n\_estimators** - The number of trees in the forest.

**min\_weight\_fraction\_leaf** - The minimum number of samples required to be at a leaf node.

**max\_features** - The number of features to consider when looking for the best

**criterion** - The function to measure the quality of a split. Supported criteria are “gini” for the Gini impurity and “entropy” for the information gain.

Parameter	Type	Conclusion
max_features	string	Contains the options ”auto”, ”sqrt” and ”log2”. From (Figure ??) we can see that ”sqrt” has a higher accuracy overall, the accuracy of ”log2” varies between the lower end and the median accuracy value.
min_samples_split	int or float	The minimum samples required to split a node has very little impact on accuracy.
criterion	string	Gives a perfect negative correlation with respect to accuracy. Correlation values being [Gini = -0.404 ], [Entropy = 0.404 ].
n_estimators	int	This is defined as the number of trees in the forest, it seems to have very little correlation but high importance.
min_samples_leaf	int or float	Gives a strong negative correlation in terms of accuracy, meaning the higher minimum samples at a leaf node, the lower the accuracy.
min_weight_fraction_leaf	float	Has a somewhat positive correlation to accuracy. e.g. total weight required at a leaf node varies between 76% and 89% accuracy

See Parameter Importance (Figure C.1)

## 5 Variation in performance with varying learning parameters in Neural Networks

### 5.1 Linear Classifier

For a linear classifier we have decided to use logistic regression. Shown below is the hyperparameter importance concerning the performance of Logistic Regression.

The accuracy fluctuates depending on which hyperparameters are used. Logistic Regression has a ‘solver’ hyperparameter in sklearn which is the algorithm to use in the optimization. We have tested the following: ‘newton-cg’, ‘lbfgs’, ‘liblinear’, ‘sag’, ‘saga’. Sag and Saga give the best results with 89% accuracy. Liblinear isn’t far behind with 88%. Other solvers have a negative impact on the performance, this can be seen in the Solver parameter importance table where their correlation values are negative.

Additionally ‘sag’ and ‘saga’ guarantee fast convergence on features with approximately the same scale [R41] which is the case with our datasets as pixel values have the same scale. This can be seen as the number of iterations between 200-800 don’t change much in the accuracy and from the scatter chart of the accuracy of the Logistic Regression vs the number of sweeps created.

Parameter	Type	Conclusion
C	float	The smaller this value is the stronger the regularization. For this we settled on 0.08027. Although this parameter did not make a big difference and provided similar results for values between 0.8 and 1.2
fit_intercept	bool	From testing we have found out that for this dataset adding a bias works better than not having bias thus True
max_iter	int	We settled on 342 maximum number of iterations taken for the solver to converge. This parameter does not have a big importance but we found 342 gives accurate results.
solver	string	Contains the options ‘newton-cg’, ‘lbfgs’, ‘liblinear’, ‘sag’, ‘saga’, default=’lbfgs’. ‘Saga’ gives the best accuracy overall. Solver is the most influential hyperparameter as we can see from the importance and correlation table
tol	float	tolerance for stopping criteria which tells the algorithm to stop searching when some tolerance is achieved. This parameter did not make a difference, we settled on 0.0002277

See Parameter Importance (Figure D.1)

## 5.2 Multilayer Perceptron

**ADD PARAM DEFINITIONS**

Parameter	Type	Conclusion
alpha	float	This has a positive correlation to accuracy as higher alpha value equates to higher accuracy.
solver	string	'Lbfgs' solver did provide the best accuracy consistently when compared to other solvers such as 'adam' and 'sgd' which supports the statement from [X002] staying that 'Lbfgs' works best for small datasets.
max_iter	int	The maximum number of iterations - In general, higher accuracy can be achieved with a larger amount of max iterations.
activation	string	Out of the four activation functions (relu, tanh, identity and logistic), relu is the only one with a positive correlation, giving the highest accuracy overall.
learning_rate	string	'adaptive' achieves the highest accuracy while, 'constant' and 'invscale' vary widely.
hidden_layer_sizes	tuple	Has a negative correlation - the number neurons in the n-th hidden layer has no effect on accuracy.

See Parameter Importance (Figure E.1)

## 5.3 Convolutional Neural Networks

### 5.3.1 Critical limitation of CNN experimental methodology:

We were unable to disable memoization between experiments over each loop, so all the parameters are restored between runs within a given experiment, consequently we noted only significant changes to the loss when training on all classes, after which convergence was achieved in a single epoch. We attempted to clear the model parameters, clear the GPU memory, and model save the randomly initialised parameters and load them back in between experiments but could not prevent the CNN model using data encoded from previous experiments. This does invalidate our results, when compared to the discrete experiments in the previous sections, however this can be a demonstrable advantage of a CNN, in this situation it consistently correctly classifies the images despite changing the dimensions of the output layer and the classification targets (albeit on the same dataset).

During all experiments with training and testing data, at worst our CNN scored an accuracy of 97% on the test dataset when classifying all classes, the CNN was capable of correct classification even when 9000 instances were moved from the training set to the testing set.

The CNN was unsurprisingly the best performing classifier, In many of the kfold experiments the CNN managed an accuracy of 100%, but this could be due in part to the memoization effect between experiments, this is particularly noticeable on the kfold experiments.

## 6 Variation in performance according to different metrics (TP Rate, FP Rate, Precision, Recall, F Measure, ROC Area)

# Appendices

## A Appendix A

### A.1 Workload split

Team member	Involvement
Lewis Wilson	text here
Chun Man	Contributed to the experimenting and analysis of Random Forest, Report writing.
Sam Fay-Hunt	text here
Kamil Szymczak	text here

As a team we are happy with everyone's contributions to the project. All team members were punctual and showed up to all scheduled meetings. Sam took the lead as project manager throughout the project delegating the workload and providing support to others.

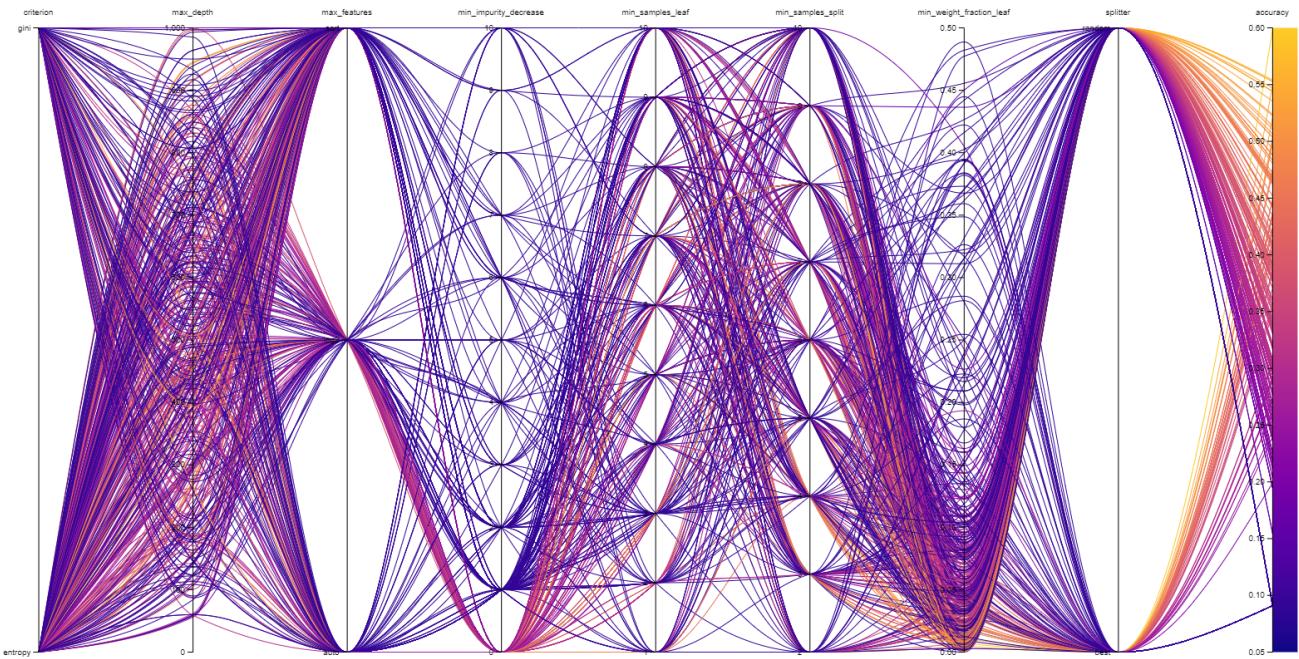
## B J48

### B.1 J48 Parameter Importance

Parameter Config	Importance	Correlation
min_weight_fraction_leaf	0.668	-0.666
min_impurity_decrease	0.052	-0.241
min_samples_leaf	0.016	0.007
max_depth	0.012	0.084
min_samples_split	0.009	-0.108

Parameter Config	Importance	Correlation
max_features.value_log2	0.133	-0.372
max_features.value_sqrt	0.003	0.196
splitter.value_best	0.049	0.051
splitter.value_random	0.043	-0.051
criterion.value_entropy	0.002	-0.022
criterion.value_gini	0.002	0.022

Figure 5: J48 Parameters



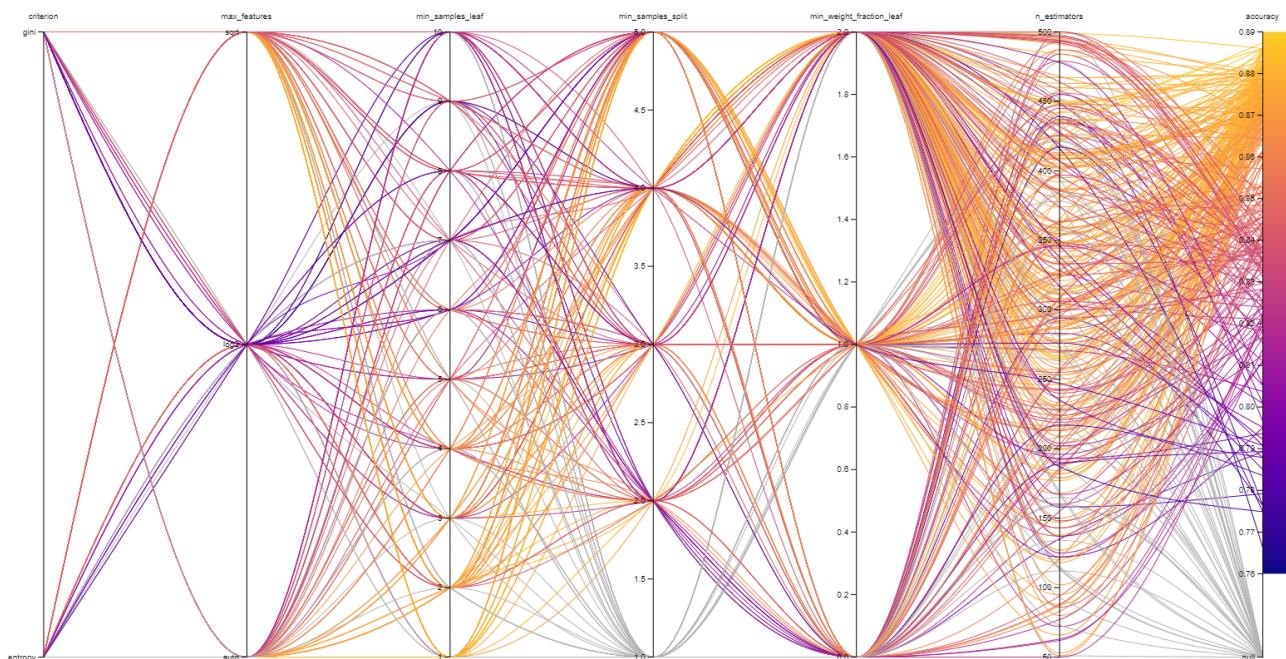
## C Random Forest

### C.1 Random Forest Parameter Importance

Parameter Config	Importance	Correlation
min_samples_split	0.005	0.306
min_samples_leaf	0.375	-0.725
n_estimators	0.016	0.092
min_weight_fraction_leaf	0.013	0.123

Parameter Config	Importance	Correlation
max_features.value_sqrt	0.001	0.563
max_features.value_log2	0.565	-0.752
criterion.value_entropy	0.012	0.404
criterion.value_gini	0.012	-0.404

Figure 6: Random Forest Parameters



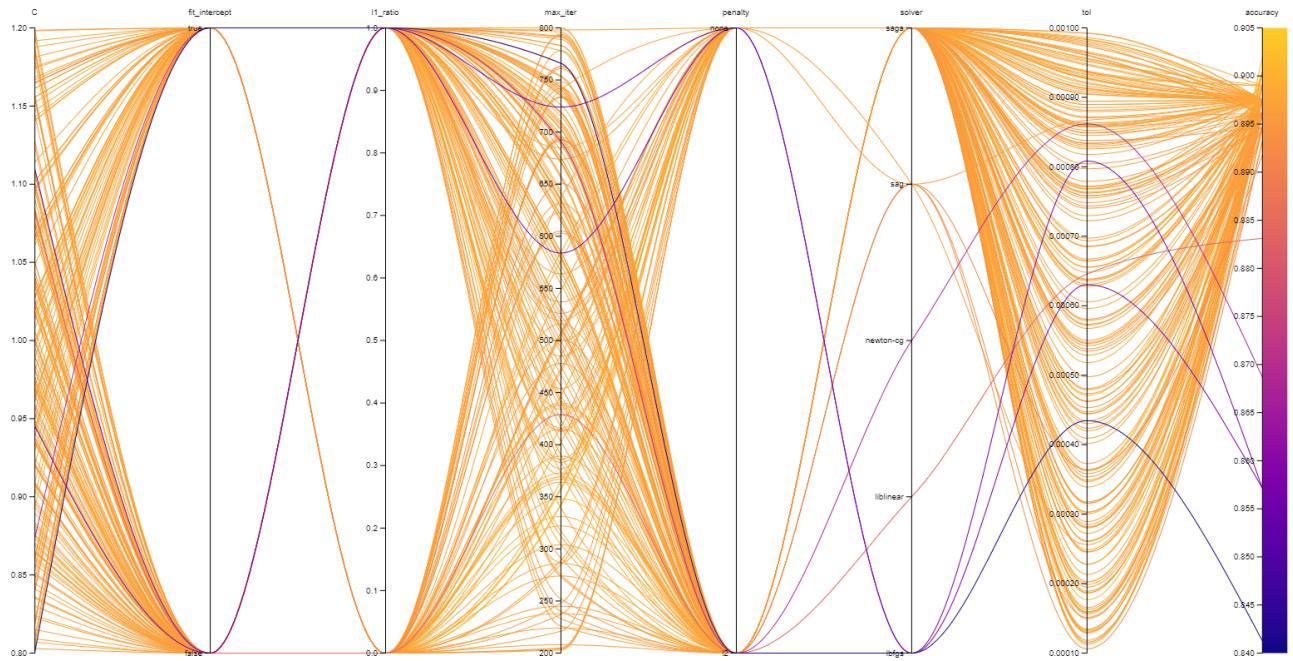
## D Linear Classifier

### D.1 Linear Classifier Parameter Importance

Solver Parameter Config	Importance	Correlation
lbfgs	0.786	-0.869
newton-cg	0.159	-0.271
liblinear	0.042	-0.037
saga	0.009	0.691
sag	0.003	0.170

Config Parameter Config	Importance	Correlation
max_iter	0.194	-0.634
l1_ratio	0.124	-0.510
fit_intercept	0.058	0.367
tol	0.046	-0.174
C	0.031	0.260

Figure 7: Random Forest Parameters



## E Multilayer Perceptron

### E.1 Multilayer Perceptron Parameter Importance

Parameter Config	Importance	Correlation
hidden_layer_sizes	0.095	-0.101
max_iter	0.091	0.072
alpha	0.061	0.171

Parameter Config	Importance	Correlation
solver.value_lbgfs	0.530	0.728
solver.value_adam	0.027	-0.245
solver.value_sgd	0.024	-0.640
activation.value_identity	0.098	-0.169
activation.value_relu	0.033	0.301
activation.value_tanh	0.018	-0.035
activation.value_logistic	0.006	-0.270
learning_rate.value_adaptive	0.012	0.507
learning_rate.value_constant	0.004	-0.442
learning_rate.value_invscaling	0.001	-0.224

Figure 8: Multilayer Perceptron Parameters

